

EC-252: COMPUTER ARCHITECTURE AND MICROPROCESSORS

Vaskar Raychoudhury

Indian Institute of Technology Roorkee



Major Advances in Computers(1)

2

- The family concept
 - ▣ IBM System/360 1964
 - ▣ DEC PDP-8
 - ▣ Separates architecture from implementation
- Microporgrammed control unit
 - ▣ Idea by Wilkes 1951
 - ▣ Produced by IBM S/360 1964
- Cache memory
 - ▣ IBM S/360 model 85 1969

Major Advances in Computers(2)

3

- Microprocessors
 - ▣ Intel 4004 1971
- Pipelining
 - ▣ Introduces parallelism into fetch execute cycle
- Multiple processors

The Next Step - RISC

4

- Reduced Instruction Set Computer
- Key features
 - ▣ Large number of general purpose registers
 - ▣ or use of compiler technology to optimize register use
 - ▣ Limited and simple instruction set
 - ▣ Emphasis on optimising the instruction pipeline

Comparison of processors

5

	Complex Instruction Set (CISC) Computer			Reduced Instruction Set (RISC) Computer		Superscalar		
Characteristic	IBM 370/168	VAX 11/780	Intel 80486	SPARC	MIPS R4000	PowerPC	Ultra SPARC	MIPS R10000
Year developed	1973	1978	1989	1987	1991	1993	1996	1996
Number of instructions	208	303	235	69	94	225		
Instruction size (bytes)	2–6	2–57	1–11	4	4	4	4	4
Addressing modes	4	22	11	1	1	2	1	1
Number of general- purpose registers	16	16	8	40 - 520	32	32	40 - 520	32
Control memory size (Kbits)	420	480	246	—	—	—	—	—
Cache size (KBytes)	64	64	8	32	128	16-32	32	64

Driving force for CISC

6

- Software costs far exceed hardware costs
- Increasingly complex high level languages
- Semantic gap
- Leads to:
 - ▣ Large instruction sets
 - ▣ More addressing modes
 - ▣ Hardware implementations of HLL statements
 - e.g. CASE (switch) on VAX

Intention of CISC

7

- Ease compiler writing
- Improve execution efficiency
 - ▣ Complex operations in microcode
- Support more complex HLLs

Execution Characteristics

8

- Operations performed
- Operands used
- Execution sequencing
- Studies have been done based on programs written in HLLs
- Dynamic studies are measured during the execution of the program

Operations

9

- Assignments
 - ▣ Movement of data
- Conditional statements (IF, LOOP)
 - ▣ Sequence control
- Procedure call-return is very time consuming
- Some HLL instruction lead to many machine code operations

Weighted Relative Dynamic Frequency of HLL Operations [PATT82a]

10

	Dynamic Occurrence		Machine-Instruction Weighted		Memory-Reference Weighted	
	Pascal	C	Pascal	C	Pascal	C
ASSIGN	45%	38%	13%	13%	14%	15%
LOOP	5%	3%	42%	32%	33%	26%
CALL	15%	12%	31%	33%	44%	45%
IF	29%	43%	11%	21%	7%	13%
GOTO	—	3%	—	—	—	—
OTHER	6%	1%	3%	1%	2%	1%

Operands

11

- Mainly local scalar variables
- Optimisation should concentrate on accessing local variables

	Pascal	C	Average
Integer Constant	16%	23%	20%
Scalar Variable	58%	53%	55%
Array/Structure	26%	24%	25%

Procedure Calls

12

- Very time consuming
- Depends on number of parameters passed
- Depends on level of nesting
- Most programs do not do a lot of calls followed by lots of returns
- Most variables are local
- (c.f. locality of reference)

Implications

13

- Best support is given by optimising most used and most time consuming features
- Large number of registers
 - ▣ Operand referencing
- Careful design of pipelines
 - ▣ Branch prediction etc.
- Simplified (reduced) instruction set

Large Register File

14

- Software solution
 - ▣ Require compiler to allocate registers
 - ▣ Allocate based on most used variables in a given time
 - ▣ Requires sophisticated program analysis
- Hardware solution
 - ▣ Have more registers
 - ▣ Thus more variables will be in registers

Registers for Local Variables

15

- ❑ Store local scalar variables in registers
- ❑ Reduces memory access
- ❑ Every procedure (function) call changes locality
- ❑ Parameters must be passed
- ❑ Results must be returned
- ❑ Variables from calling programs must be restored

Register Windows

16

- Only few parameters
- Limited range of depth of call
- Use multiple small sets of registers
- Calls switch to a different set of registers
- Returns switch back to a previously used set of registers

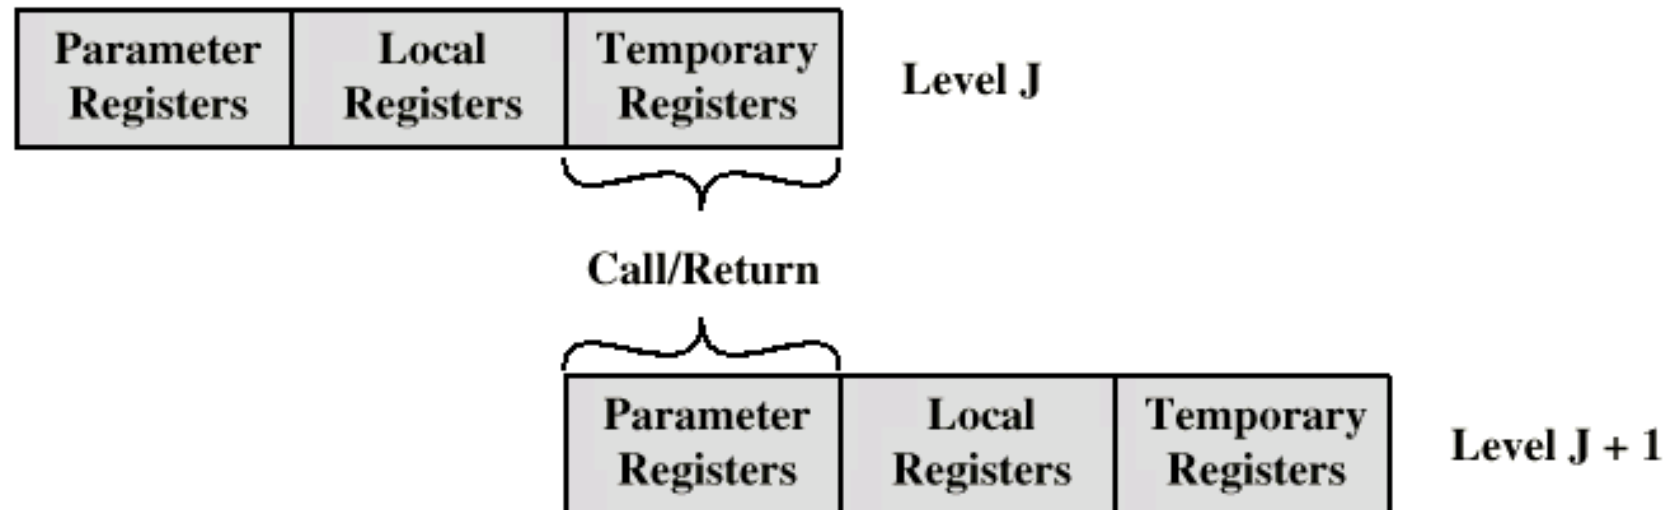
Register Windows cont.

17

- Three areas within a register set
 - ▣ Parameter registers
 - ▣ Local registers
 - ▣ Temporary registers
 - ▣ Temporary registers from one set overlap parameter registers from the next
 - ▣ This allows parameter passing without moving data


Overlapping Register Windows

18



Circular Buffer diagram

19



Operation of Circular Buffer

20

- When a call is made, a current window pointer is moved to show the currently active register window
- If all windows are in use, an interrupt is generated and the oldest window (the one furthest back in the call nesting) is saved to memory
- A saved window pointer indicates where the next saved windows should restore to

Global Variables

21

- Allocated by the compiler to memory
 - ▣ Inefficient for frequently accessed variables
- Have a set of registers for global variables

RISC Characteristics

22

- ❑ One instruction per cycle
- ❑ Register to register operations
- ❑ Few, simple addressing modes
- ❑ Few, simple instruction formats
- ❑ Hardwired design (no microcode)
- ❑ Fixed instruction format
- ❑ More compile time/effort

CISC vs RISC

23

- Hard to Distinguish Now. Boundary is getting vague.
- Academia don't Care
- Industry doesn't Care (Except for Advertisements)
- CISC
 - ▣ *Effectively realizes one particular High Level Language Computer System in HW* - recurring **HW development costs** when change needed
- RISC
 - ▣ *Allows effective realization of any High Level Language Computer System in SW* - recurring **SW development costs** when change needed

Registers v Cache (w.r.to RISC)

24

Large Register File	Cache
All local scalars	Recently-used local scalars
Individual variables	Blocks of memory
Compiler-assigned global variables	Recently-used global variables
Save/Restore based on procedure nesting depth	Save/Restore based on cache replacement algorithm
Register addressing	Memory addressing

RISC Pipelining

25

- Most instructions are register to register
- Two phases of execution
 - ▣ I: Instruction fetch
 - ▣ E: Execute
 - ALU operation with register input and output
- For load and store
 - ▣ I: Instruction fetch
 - ▣ E: Execute
 - Calculate memory address
 - ▣ D: Memory
 - Register to memory or memory to register operation

Effects of Pipelining

26

Load $rA \leftarrow M$	I	E	D								
Load $rB \leftarrow M$				I	E	D					
Add $rC \leftarrow rA + rB$							I	E			
Store $M \leftarrow rC$									I	E	D
Branch X										I	E

(a) Sequential execution

Load $rA \leftarrow M$	I	E	D								
Load $rB \leftarrow M$		I		E	D						
Add $rC \leftarrow rA + rB$				I		E					
Store $M \leftarrow rC$							I	E	D		
Branch X							I			E	
NOOP										I	E

(b) Two-stage pipelined timing

Load $rA \leftarrow M$	I	E	D								
Load $rB \leftarrow M$		I	E	D							
NOOP			I	E							
Add $rC \leftarrow rA + rB$				I	E						
Store $M \leftarrow rC$						I	E	D			
Branch X							I	E			
NOOP									I	E	

(c) Three-stage pipelined timing

Load $rA \leftarrow M$	I	E ₁	E ₂	D							
Load $rB \leftarrow M$		I	E ₁	E ₂	D						
NOOP			I	E ₁	E ₂						
NOOP				I	E ₁	E ₂					
Add $rC \leftarrow rA + rB$					I	E ₁	E ₂				
Store $M \leftarrow rC$						I	E ₁	E ₂	D		
Branch X							I	E ₁	E ₂		
NOOP								I	E ₁	E ₂	
NOOP									I	E ₁	E ₂

(d) Four-stage pipelined timing

Optimization of Pipelining

27

- Delayed branch
 - ▣ Does not take effect until after execution of following instruction
 - ▣ This following instruction is the delay slot
- Delayed Load
 - ▣ Register to be target is locked by processor
 - ▣ Continue execution of instruction stream until register required
 - ▣ Idle until load complete
 - ▣ Re-arranging instructions can allow useful work whilst loading
- Loop Unrolling
 - ▣ Replicate body of loop a number of times
 - ▣ Iterate loop fewer times
 - ▣ Reduces loop overhead
 - ▣ Increases instruction parallelism
 - ▣ Improved register, data cache or TLB locality

Normal and Delayed Branch

28

Address	Normal Branch	Delayed Branch	Optimized Delayed Branch
100	LOAD X, rA	LOAD X, rA	LOAD X, rA
101	ADD 1, rA	ADD 1, rA	JUMP 105
102	JUMP 105	JUMP 106	ADD 1, rA
103	ADD rA, rB	NOOP	ADD rA, rB
104	SUB rC, rB	ADD rA, rB	SUB rC, rB
105	STORE rA, Z	SUB rC, rB	STORE rA, Z
106		STORE rA, Z	

Use of Delayed Branch

29

- The delayed branch means that
 - ▣ the instruction following the branch is always executed before the PC is modified to perform the branch

100 LOAD X, rA
 101 ADD 1, rA
 102 JUMP 105
 103 ADD rA, rB
 105 STORE rA, Z

Time →

1	2	3	4	5	6	7	8
I	E	D					
	I		E				
			I	E			
				I			
					I	E	D

(a) Traditional Pipeline

100 LOAD X, rA
 101 ADD 1, rA
 102 JUMP 106
 103 NOOP
 106 STORE rA, Z

I	E	D				
	I	E				
		I	E			
			I	E		
				I	E	D

(b) RISC Pipeline with Inserted NOOP

100 LOAD X, rA
 101 JUMP 105
 102 ADD 1, rA
 105 STORE rA, Z

I	E	D			
	I	E			
		I	E		
			I	E	D

(c) Reversed Instructions

Delayed Branch and Delay Slot

30

- A delayed branch specifies that the
 - ▣ jump to a new location happens *after* the next instruction
 - ▣ that next instruction is the one (called, the delay slot) loaded after the branch
- Branch Delay Slot
 - ▣ Always fetch the instruction after the branch from the instruction cache, and always execute it, even if the branch is taken
- MIPS-X use a double branch delay slot
 - ▣ executes a pair of instructions following a branch instruction before the branch takes effect

Disadvantages of Delayed Branch

31

- Delayed branches have been criticized as a poor short-term choice in ISA design
 - ▣ Compilers typically have some difficulty finding logically independent instructions to place after the branch
 - ▣ Better to insert NOPs/NOOPs into the delay slots